

Deep Learning

10 Building Blocks of CNNs

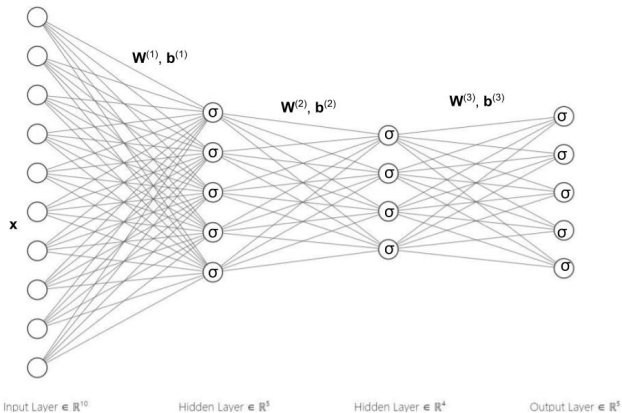
Dr. Konda Reddy Mopuri
Dept. of AI, IIT Hyderabad
Jan-May 2023

- The Convolutional Neural Networks

- The Convolutional Neural Networks
- Class of ANNs that are Shift/Space invariant
 - Makes CNNs very well suited for *Signal Processing* (Why?).


An MLP

- Input is a vector



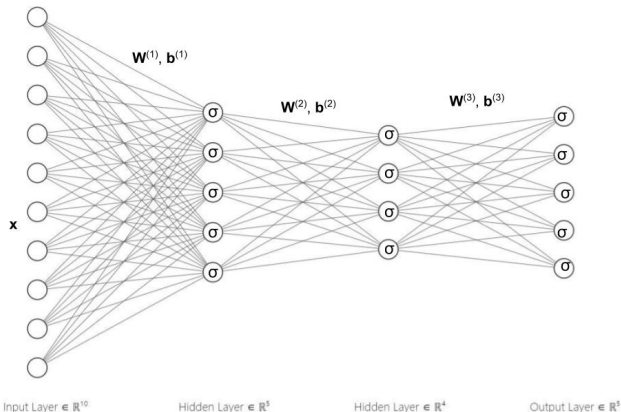
An MLP

- Input is a vector
- Series of densely connected hidden layers



An MLP

- Input is a vector
- Series of densely connected hidden layers
- **Neurons in each layer are independent!**



An MLP for processing an image

- Say, we want to process a 200×200 RGB image

An MLP for processing an image

- Say, we want to process a 200×200 RGB image
- Vectorizing leads to $200 \times 200 \times 3 \rightarrow 120K$ neurons in the input layer

An MLP for processing an image

- Say, we want to process a 200×200 RGB image
- Vectorizing leads to $200 \times 200 \times 3 \rightarrow 120K$ neurons in the input layer
- A hidden layer of same size leads to $\approx 1.44e^{10}$ weights $\rightarrow \approx 58GB :-()$

An MLP for processing an image

- Full connectivity blows the number of weights \rightarrow hardware limits, overfitting, etc.

An MLP for processing an image

- Full connectivity blows the number of weights \rightarrow hardware limits, overfitting, etc.
- Flattening removes the structure

Large Signals

- Have invariance in translation

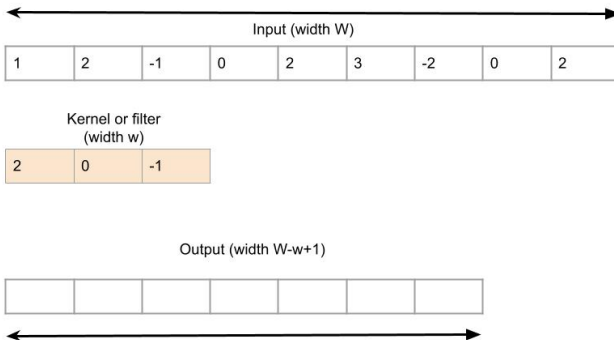
Large Signals

- Have invariance in translation
- Features may occur at different locations in the signal

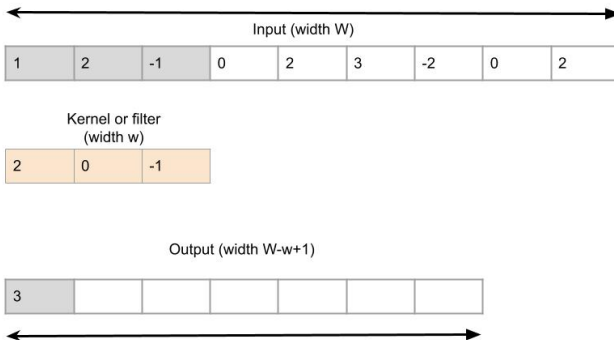
Large Signals

- Have invariance in translation
- Features may occur at different locations in the signal
- **Convolution** incorporates this idea: Applies same linear operation at all the locations and preserves the structure

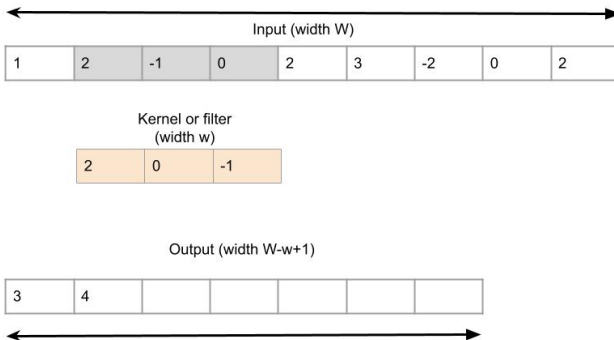
Convolution




Convolution



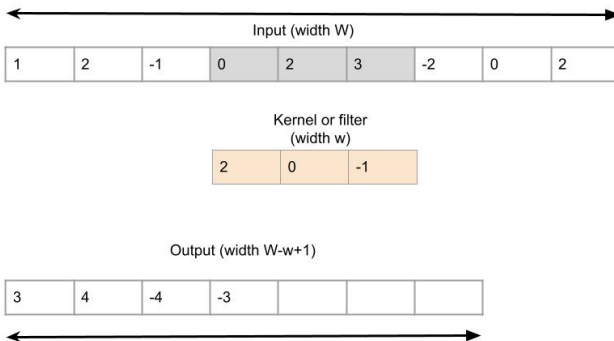
Convolution



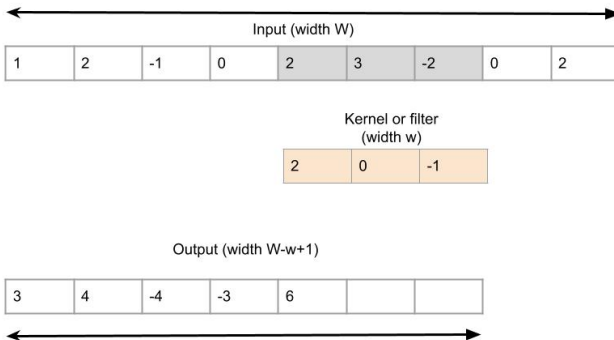
Convolution



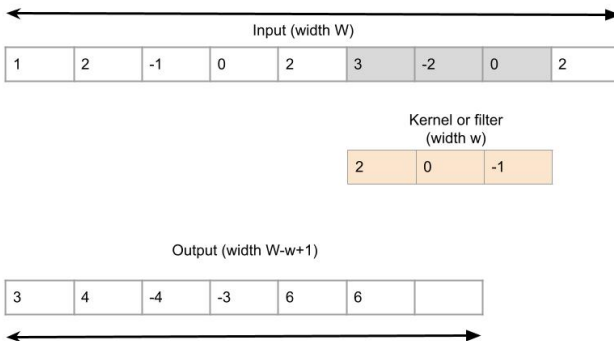
Convolution



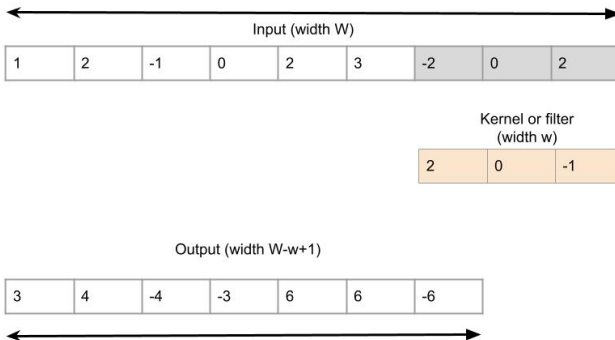
Convolution



Convolution



Convolution



Convolution

- Preserves the structure

- Preserves the structure
 - if the i/p is a 2D tensor \rightarrow o/p is also a 2D tensor

- Preserves the structure
 - if the i/p is a 2D tensor \rightarrow o/p is also a 2D tensor
 - There exist a relation between the locations of i/p and o/p values

- Let $\mathbf{x} = (x_1, x_2, \dots, x_W)$ is the input, $\mathbf{k} = (k_1, k_2, \dots, k_w)$ is the kernel

- Let $\mathbf{x} = (x_1, x_2, \dots, x_W)$ is the input, $\mathbf{k} = (k_1, k_2, \dots, k_w)$ is the kernel
- The result $(x \circledast k)$ of convolving \mathbf{x} with \mathbf{k} will be a 1D tensor of size $W - w + 1$

$$\begin{aligned}(x \circledast k)_i &= \sum_{j=1}^w x_{i-1+j} k_j \\ &= (x_i, \dots, x_{i+w-1}) \cdot \mathbf{k}\end{aligned}$$

Convolution

- Powerful feature extractor

Convolution


- Powerful feature extractor
- For instance, it can perform differential operation and look for interesting patterns in the input

Convolution

- Powerful feature extractor
- For instance, it can perform differential operation and look for interesting patterns in the input



$$(0, 0, 0, 1, 2, 3, 4, 4, 4, 4) \otimes (-1, 1) = (0, 0, 1, 1, 1, 1, 0, 0, 0)$$




Convolution

- Powerful feature extractor
- For instance, it can perform differential operation and look for interesting patterns in the input



$$(0, 0, 1, 1, 0, 0.1, 0.2, 1, 1, 1, 0) \circledast (1, 1) = (0, 1, 2, 1, 0.1, 0.3, 1.2, 2, 2, 1)$$





- Naturally generalizes to multiple dimensions

- Naturally generalizes to multiple dimensions
- CNNs process 3D tensors of size $C \times H \times W$ with kernels of size $C \times h \times w$ and result in 2D tensors of size $H - h + 1 \times W - w + 1$


- Naturally generalizes to multiple dimensions
- CNNs process 3D tensors of size $C \times H \times W$ with kernels of size $C \times h \times w$ and result in 2D tensors of size $H - h + 1 \times W - w + 1$
- Note that we generally refer to these inputs as 2D signal (despite having C channels) (Why?)

2D Convolution




kernel


2D Convolution




2D Convolution




2D Convolution




2D Convolution




2D Convolution




2D Convolution




2D Convolution




2D Convolution



2D Convolution



2D Convolution




2D Convolution

- Kernel is not convolved in the channel dimension


2D Convolution

- Kernel is not convolved in the channel dimension
- Another way to interpret convolution is that an affine function is applied on an input block of size $C \times h \times w$



2D Convolution


- Kernel is not convolved in the channel dimension
- Another way to interpret convolution is that an affine function is applied on an input block of size $C \times h \times w$



- Same affine function is applied on all such blocks in the input

2D Convolution

- Kernel is not convolved in the channel dimension
- Another way to interpret convolution is that an affine function is applied on an input block of size $C \times h \times w$



- Same affine function is applied on all such blocks in the input

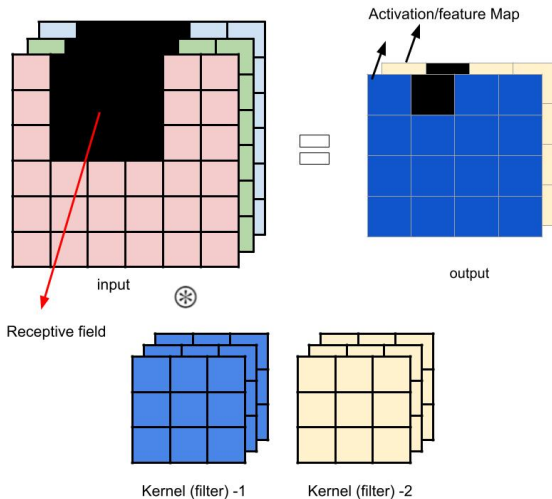
- Preserves the input structure

- Preserves the input structure
 - 1D signal outputs 1D signal, 2D signal outputs 2D signal

- Preserves the input structure
 - 1D signal outputs 1D signal, 2D signal outputs 2D signal
 - Adjacent components in o/p are influenced by adjacent parts in the i/p

- Preserves the input structure
 - 1D signal outputs 1D signal, 2D signal outputs 2D signal
 - Adjacent components in o/p are influenced by adjacent parts in the i/p
- If the channel dimension has a metric meaning (e.g. time) 3D convolution can be employed (e.g. frames in a video)

Terminology in Convolution



Convolution function in PyTorch

- `F.conv2d(input, weight, bias=None, stride=1, padding=0, dilation=1, groups=1)`

Convolution function in PyTorch

- `F.conv2d(input, weight, bias=None, stride=1, padding=0, dilation=1, groups=1)`
- `weight` is $D \times C \times h \times w$ dimensional kernels

Convolution function in PyTorch

- `F.conv2d(input, weight, bias=None, stride=1, padding=0, dilation=1, groups=1)`
- `weight` is $D \times C \times h \times w$ dimensional kernels
- `bias` D dimensional

Convolution function in PyTorch

- `F.conv2d(input, weight, bias=None, stride=1, padding=0, dilation=1, groups=1)`
- `weight` is $D \times C \times h \times w$ dimensional kernels
- `bias` D dimensional
- `input` is $N \times C \times H \times W$ dimensional signal

Convolution function in PyTorch

- `F.conv2d(input, weight, bias=None, stride=1, padding=0, dilation=1, groups=1)`
- `weight` is $D \times C \times h \times w$ dimensional kernels
- `bias` D dimensional
- `input` is $N \times C \times H \times W$ dimensional signal
- Output is $N \times D \times (H - h + 1) \times (W - w + 1)$ tensor

Convolution function in PyTorch

- `F.conv2d(input, weight, bias=None, stride=1, padding=0, dilation=1, groups=1)`
- `weight` is $D \times C \times h \times w$ dimensional kernels
- `bias` D dimensional
- `input` is $N \times C \times H \times W$ dimensional signal
- Output is $N \times D \times (H - h + 1) \times (W - w + 1)$ tensor
- Autograd compliant

Convolution function in PyTorch

```
input = torch.empty(128, 3, 20, 20).normal_()
weight = torch.empty(5, 3, 5, 5).normal_()
bias = torch.empty(5).normal_()
output = F.conv2d(input, weight, bias)
output.size()
torch.Size([128, 5, 16, 16])
```

Look/Access the filters

```
weight[0,0]  
tensor([[ -0.6974,  0.1342, -0.2632, -0.4672,  0.1827],  
        [ -0.1184, -0.2164,  0.2772, -0.1099,  0.0103],  
        [ -0.8272,  0.3580,  0.2398, -0.5795, -0.9472],  
        [ -1.1734, -0.1019,  0.7394,  0.3342,  0.1699],  
        [  1.9271,  0.1250,  0.4222,  0.2014,  1.1100]])
```

Conv layer in PyTorch

- Class `torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True)`

Conv layer in PyTorch

- Class `torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True)`
- `kernel_size` can be either a pair (h , w) or a single value k interpreted as (k , k).

Conv layer in PyTorch

- Class `torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True)`
- `kernel_size` can be either a pair (h , w) or a single value k interpreted as (k , k).
- Encloses the convolution as a module

Conv layer in PyTorch

- Class `torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True)`
- `kernel_size` can be either a pair (h , w) or a single value k interpreted as (k , k).
- Encloses the convolution as a module
- Initializes the kernel parameters and biases as random

Conv layer in PyTorch

```
f = nn.Conv2d(in_channels = 3, out_channels = 5,  
kernel_size = (2, 3))  
for n, p in f.named_parameters():  
...print(n, p.size())  
  
>>weight torch.Size([5, 3, 2, 3])  
>>bias torch.Size([5])
```

Conv layer in PyTorch

```
f = nn.Conv2d(in_channels = 3, out_channels = 5,  
kernel_size = (2, 3))  
for n, p in f.named_parameters():  
    ...print(n, p.size())  
  
>>weight torch.Size([5, 3, 2, 3])  
>>bias torch.Size([5])  
  
input = torch.empty(128, 3, 28, 28).normal_()  
output = f(input)  
output.size()  
>>torch.Size([128, 5, 27, 26])
```

Padding in Convolution

- Adds zeros around the input


Padding in Convolution

- Adds zeros around the input
- Takes care of size reduction after convolution


Padding in Convolution

- Adds zeros around the input
- Takes care of size reduction after convolution
- Instead of zeros, one may pad with signal values at the edges

Padding in Convolution



Padding in Convolution



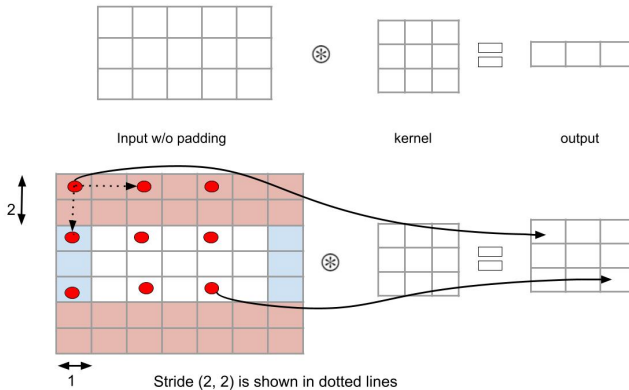
Stride in Convolution

- Specifies the step size taken while performing convolution

Stride in Convolution

- Specifies the step size taken while performing convolution
- Default value is 1, i.e., move the kernel across the signal densely (without skipping)

Padding and Stride in Convolution



Dilation in Convolution

- Manipulates the size of the kernel via expanding its size without adding weights.

Dilation in Convolution

- Manipulates the size of the kernel via expanding its size without adding weights.
- In other words, it inserts 0s in between the kernel values


Output size of the Convolution

- Input width - W , Kernel size - k , Padding - p , and stride - s


Output size of the Convolution

- Input width - W , Kernel size - k , Padding - p , and stride - s
- Output width = $\frac{W-k+2p}{s} + 1$ (similarly for the height)

Without Dilation



Dilation (2, 2)



- Expands the kernel by adding rows and columns of zeros

- Expands the kernel by adding rows and columns of zeros
- Default value for dilation is 1, i.e., no zeros placed

- Expands the kernel by adding rows and columns of zeros
- Default value for dilation is 1, i.e., no zeros placed
- Any higher value of dilation makes the kernel sparse

- Expands the kernel by adding rows and columns of zeros
- Default value for dilation is 1, i.e., no zeros placed
- Any higher value of dilation makes the kernel sparse
- Dilation increases the receptive field

- Expands the kernel by adding rows and columns of zeros
- Default value for dilation is 1, i.e., no zeros placed
- Any higher value of dilation makes the kernel sparse
- Dilation increases the receptive field
- It is referred to as 'atrous' convolution

Pooling

- Groups multiple activations and replaces by a representative one

- Groups multiple activations and replaces by a representative one
- Reduces the dimensionality of the signal progressively → considers non-overlapping stride

- Groups multiple activations and replaces by a representative one
- Reduces the dimensionality of the signal progressively → considers non-overlapping stride
- Also called sub-sampling layer


- Groups multiple activations and replaces by a representative one
- Reduces the dimensionality of the signal progressively → considers non-overlapping stride
- Also called sub-sampling layer
- Generally found between two convolution layers (and parameter free)

Max Pooling

- Standard in CNNs


Max Pooling

- Standard in CNNs
- Computes maximum value over a non-overlapping blocks in the input



Average Pooling

- Computes the average of the receptive field



Pooling in 2D

- Same as 1D, but the receptive field is 2D and non-overlapping






Figure credits: Preston Hoang and Quora

Pooling in 2D

- Contrary to Convolution, Pooling applies channel wise

Pooling in 2D

- Contrary to Convolution, Pooling applies channel wise
- No reduction in number of channels, only spatial size reduction



Pooling provides weak invariance

- Operation is invariant to any permutation within the block

Pooling provides weak invariance

- Operation is invariant to any permutation within the block
- Withstands deformations caused by local translations

Max_Pooling PyTorch

```
F.max_pool2d(input, kernel_size, stride=None, padding=0,  
dilation=1, ceil_mode=False, return_indices=False)
```

- Applies max pooling on each of the channels separately

Max_Pooling PyTorch

```
F.max_pool2d(input, kernel_size, stride=None, padding=0,  
dilation=1, ceil_mode=False, return_indices=False)
```

- Applies max pooling on each of the channels separately
- input is $N \times C \times H \times W$ tensor

Max_Pooling PyTorch

```
F.max_pool2d(input, kernel_size, stride=None, padding=0,  
dilation=1, ceil_mode=False, return_indices=False)
```

- Applies max pooling on each of the channels separately
- `input` is $N \times C \times H \times W$ tensor
- `kernel_size` is (h, w) or k

Max_Pooling PyTorch

```
F.max_pool2d(input, kernel_size, stride=None, padding=0,  
dilation=1, ceil_mode=False, return_indices=False)
```

- Applies max pooling on each of the channels separately
- input is $N \times C \times H \times W$ tensor
- kernel_size is (h, w) or k
- Result would be a tensor of size $N \times C \times \lfloor H/h \rfloor \times \lfloor W/w \rfloor$

Pooling in PyTorch

- Default stride is the kernel size (for convolution, it is 1)

Pooling in PyTorch

- Default stride is the kernel size (for convolution, it is 1)
- But, it can be modulated if required

Pooling in PyTorch

- Default stride is the kernel size (for convolution, it is 1)
- But, it can be modulated if required
- Default padding is zero

Pooling Layer in PyTorch

```
class torch.nn.MaxPool2d(kernel_size, stride=None,  
padding=0, dilation=1, return_indices=False,  
ceil_mode=False)
```

Putting it all together

Architecture of a simple CNN





Figure credits: Adit Deshpande


Architecture of a simple CNN



- Initially Conv layer with nonlinearity

Figure credits: Adit Deshpande


Architecture of a simple CNN



- Initially Conv layer with nonlinearity
- Followed by a few Conv + Nonlinearity layers

Figure credits: Adit Deshpande


Architecture of a simple CNN



- Initially Conv layer with nonlinearity
- Followed by a few Conv + Nonlinearity layers
- Have Pooling layers in between Conv layers → reduce the feature map size sufficiently

Figure credits: Adit Deshpande


Architecture of a simple CNN



- Initially Conv layer with nonlinearity
- Followed by a few Conv + Nonlinearity layers
- Have Pooling layers in between Conv layers → reduce the feature map size sufficiently
- Vectorize and fully connected layers

Figure credits: Adit Deshpande


Architecture of a simple CNN



INPUT \rightarrow $[\text{CONV} \rightarrow \text{RELU}] * N \rightarrow \text{POOL} * M \rightarrow [\text{FC} \rightarrow \text{RELU}] * K \rightarrow \text{FC}$

Figure credits: Adit Deshpande


Architecture of a simple CNN



INPUT \rightarrow $[[\text{CONV} \rightarrow \text{RELU}] * N \rightarrow \text{POOL}] * M \rightarrow [\text{FC} \rightarrow \text{RELU}] * K \rightarrow$
 FC

Figure credits: Adit Deshpande

Architecture of a simple CNN



INPUT \rightarrow $[[\text{CONV} \rightarrow \text{RELU}] * N \rightarrow \text{POOL}] * M \rightarrow [\text{FC} \rightarrow \text{RELU}] * K \rightarrow$
 FC

Figure credits: Adit Deshpande

Case study: LeNet-like architecture

input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>			

Case study: LeNet-like architecture

input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$		

Case study: LeNet-like architecture

input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32 \cdot (5^2 + 1)$ $= 832$	

Case study: LeNet-like architecture

input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32 \cdot (5^2 + 1)$ $= 832$	$32 \cdot 24^2 \cdot 5^2$ $= 460800$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>			

Case study: LeNet-like architecture

input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32 \cdot (5^2 + 1)$ $= 832$	$32 \cdot 24^2 \cdot 5^2$ $= 460800$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	0	0

Case study: LeNet-like architecture

input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32 \cdot (5^2 + 1)$ $= 832$	$32 \cdot 24^2 \cdot 5^2$ $= 460800$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0

Case study: LeNet-like architecture

input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32.(5^2 + 1)$ $= 832$	$32.24^2.5^2$ $= 460800$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ <code>nn.conv2d(32, 64, kernel_size=5)</code>			

Case study: LeNet-like architecture

input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32.(5^2 + 1)$ $= 832$	$32.24^2.5^2$ $= 460800$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ <code>nn.conv2d(32, 64, kernel_size=5)</code>	$64 \times 4 \times 4$		

Case study: LeNet-like architecture

input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32.(5^2 + 1)$ $= 832$	$32.24^2.5^2$ $= 460800$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ <code>nn.conv2d(32, 64, kernel_size=5)</code>	$64 \times 4 \times 4$	$64.(32.5^2 + 1)$ $= 51264$	

Case study: LeNet-like architecture

input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32.(5^2 + 1)$ $= 832$	$32.24^2.5^2$ $= 460800$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ <code>nn.conv2d(32, 64, kernel_size=5)</code>	$64 \times 4 \times 4$	$64.(32.5^2 + 1)$ $= 51264$	$64.32.4^2.5^2$ $= 819200$

Case study: LeNet-like architecture

input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32.(5^2 + 1)$ $= 832$	$32.24^2.5^2$ $= 460800$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ <code>nn.conv2d(32, 64, kernel_size=5)</code>	$64 \times 4 \times 4$	$64.(32.5^2 + 1)$ $= 51264$	$64.32.4^2.5^2$ $= 819200$
$64 \times 4 \times 4$ <code>F.max_pool2d(., kernel_size=2)</code>	$64 \times 2 \times 2$	0	0

Case study: LeNet-like architecture

input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32.(5^2 + 1)$ $= 832$	$32.24^2.5^2$ $= 460800$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ <code>nn.conv2d(32, 64, kernel_size=5)</code>	$64 \times 4 \times 4$	$64.(32.5^2 + 1)$ $= 51264$	$64.32.4^2.5^2$ $= 819200$
$64 \times 4 \times 4$ <code>F.max_pool2d(., kernel_size=2)</code>	$64 \times 2 \times 2$	0	0
$64 \times 2 \times 2$ / <code>F.relu(.)</code>	$64 \times 2 \times 2$	0	0

Case study: LeNet-like architecture

input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32.(5^2 + 1)$ $= 832$	$32.24^2.5^2$ $= 460800$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ <code>nn.conv2d(32, 64, kernel_size=5)</code>	$64 \times 4 \times 4$	$64.(32.5^2 + 1)$ $= 51264$	$64.32.4^2.5^2$ $= 819200$
$64 \times 4 \times 4$ <code>F.max_pool2d(., kernel_size=2)</code>	$64 \times 2 \times 2$	0	0
$64 \times 2 \times 2$ / <code>F.relu(.)</code>	$64 \times 2 \times 2$	0	0
$64 \times 2 \times 2$ <code>x.view(-1,256)</code>	256	0	0
256 <code>nn.Linear(256,200)</code>	200		

Case study: LeNet-like architecture

input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32.(5^2 + 1)$ $= 832$	$32.24^2.5^2$ $= 460800$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ <code>nn.conv2d(32, 64, kernel_size=5)</code>	$64 \times 4 \times 4$	$64.(32.5^2 + 1)$ $= 51264$	$64.32.4^2.5^2$ $= 819200$
$64 \times 4 \times 4$ <code>F.max_pool2d(., kernel_size=2)</code>	$64 \times 2 \times 2$	0	0
$64 \times 2 \times 2$ / <code>F.relu(.)</code>	$64 \times 2 \times 2$	0	0
$64 \times 2 \times 2$ <code>x.view(-1,256)</code>	256	0	0
256 <code>nn.Linear(256,200)</code>	200	$200(256+1)=51400$	$200.256=51200$

Case study: LeNet-like architecture

input size/ layer information	output size	# parameters	# products
$1 \times 28 \times 28$ <code>nn.Conv2d(1, 32, kernel_size=5)</code>	$32 \times 24 \times 24$	$32.(5^2 + 1)$ $= 832$	$32.24^2.5^2$ $= 460800$
$32 \times 24 \times 24$ <code>F.max_pool2d(., kernel_size=3)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ / <code>F.relu(.)</code>	$32 \times 8 \times 8$	0	0
$32 \times 8 \times 8$ <code>nn.conv2d(32, 64, kernel_size=5)</code>	$64 \times 4 \times 4$	$64.(32.5^2 + 1)$ $= 51264$	$64.32.4^2.5^2$ $= 819200$
$64 \times 4 \times 4$ <code>F.max_pool2d(., kernel_size=2)</code>	$64 \times 2 \times 2$	0	0
$64 \times 2 \times 2$ / <code>F.relu(.)</code>	$64 \times 2 \times 2$	0	0
$64 \times 2 \times 2$ <code>x.view(-1,256)</code>	256	0	0
256 <code>nn.Linear(256,200)</code>	200	0	0
200 / <code>F.relu(.)</code>	200	$200(256+1)=51400$	$200.256=51200$
200 <code>nn.Linear(200,10)</code>	10	0	0
		$10(200+1)=2010$	$10.200=2000$

Recent architectures are far more sophisticated

- Note that LeNet is a classical architecture and does not reflect the recent CNNs in complexity

Recent architectures are far more sophisticated

- Note that LeNet is a classical architecture and does not reflect the recent CNNs in complexity
- Recent CNN architectures are far more sophisticated [Contents of the next lecture(s)]
 - More depth
 - Machinery to handle the depth